

# The Java Fork-Join Pool Framework

## (Part 3)

**Douglas C. Schmidt**

**[d.schmidt@vanderbilt.edu](mailto:d.schmidt@vanderbilt.edu)**

**[www.dre.vanderbilt.edu/~schmidt](http://www.dre.vanderbilt.edu/~schmidt)**



**Professor of Computer Science**


**Institute for Software  
Integrated Systems**

**Vanderbilt University  
Nashville, Tennessee, USA**



# Learning Objectives in this Part of the Lesson

- Understand how the Java fork-join framework processes tasks in parallel
- Recognize the structure & functionality of the fork-join framework
- Know how the fork-join framework is implemented internally




---

# Java Fork-Join Pool Framework Internals

# Java Fork-Join Pool Framework Internals

- Each worker thread in a fork-join pool runs a loop that scans for (sub-)tasks to execute




# Java Fork-Join Pool Framework Internals

- Each worker thread in a fork-join pool runs a loop that scans for (sub-)tasks to execute
  - The goal is to keep the worker threads as busy as possible!



# Java Fork-Join Pool Framework Internals

- Each worker thread in a fork-join pool runs a loop that scans for (sub-)tasks to execute
  - The goal is to keep the worker threads as busy as possible!
  - A worker thread only blocks waiting for work if no (sub-)tasks are available to run





# Java Fork-Join Pool Framework Internals


- Each worker thread in a fork-join pool runs a loop that scans for (sub-)tasks to execute
  - The goal is to keep the worker threads as busy as possible!
- A worker thread only blocks waiting for work if no (sub-)tasks are available to run



Blocking a working thread is very costly on modern processors

# Java Fork-Join Pool Framework Internals


- Each worker thread in a fork-join pool runs a loop that scans for (sub-)tasks to execute
  - The goal is to keep the worker threads as busy as possible!
  - A worker thread only blocks waiting for work if no (sub-)tasks are available to run
- Each worker thread therefore checks multiple input sources for (sub-)tasks to execute





# Java Fork-Join Pool Framework Internals

- A worker thread has a “double-ended queue” (aka “deque”) that serves as its main input source



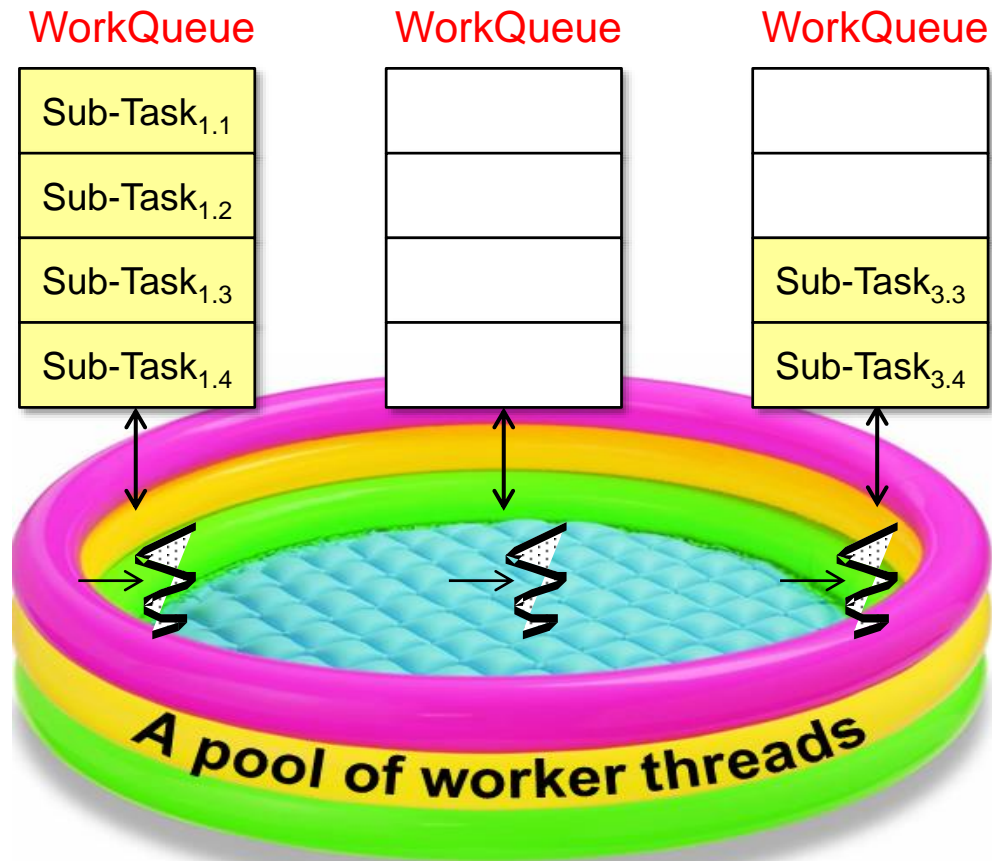
See [en.wikipedia.org/wiki/Double-ended\\_queue](https://en.wikipedia.org/wiki/Double-ended_queue)

# Java Fork-Join Pool Framework Internals

- A worker thread has a “double-ended queue” (aka “deque”) that serves as its main input source
- Implemented by WorkQueue

```
<<Java Class>>
WorkQueue

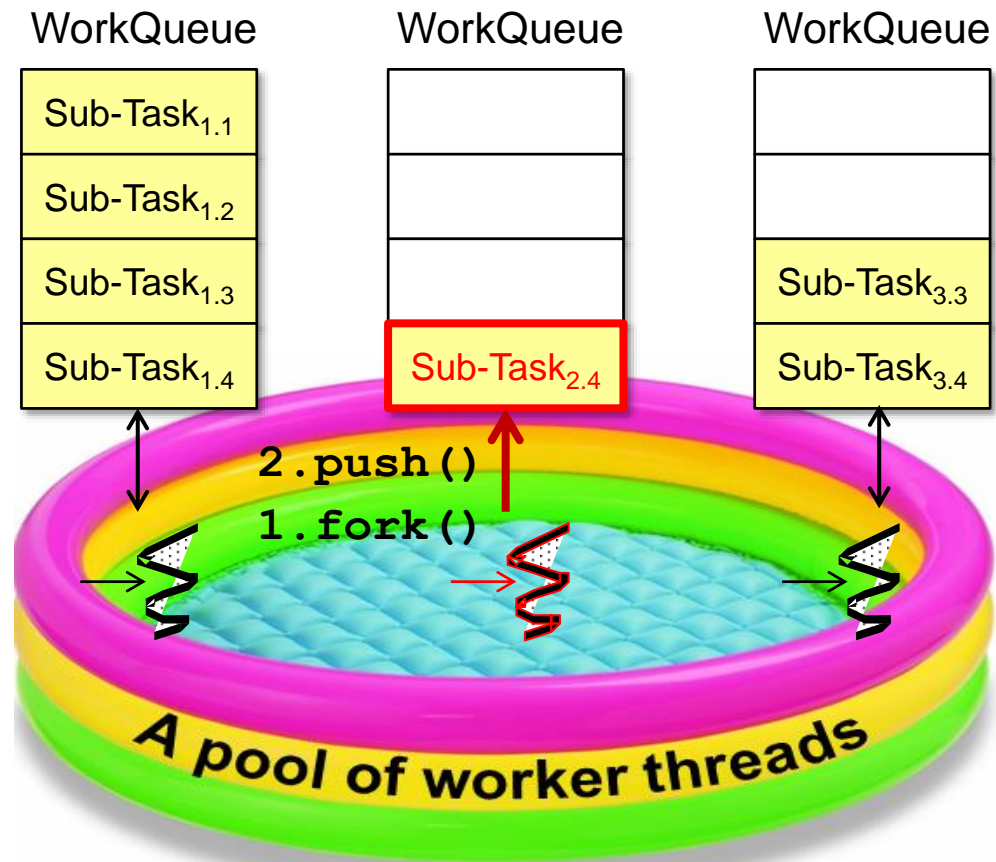
WorkQueue(ForkJoinPool,ForkJoinWorkerThread,int,int)
queueSize():int
isEmpty():boolean
push(ForkJoinTask<?>):void
growArray():ForkJoinTask<?>
pop():ForkJoinTask<?>
pollAt(int):ForkJoinTask<?>
poll():ForkJoinTask<?>
peek():ForkJoinTask<?>
cancelAll():void
pollAndExecAll():void
runTask(ForkJoinTask<?>):void
tryRemoveAndExec(ForkJoinTask<?>):boolean
isApparentlyUnblocked():boolean
```



See [java8/util/concurrent/ForkJoinPool.java](http://java8/util/concurrent/ForkJoinPool.java)

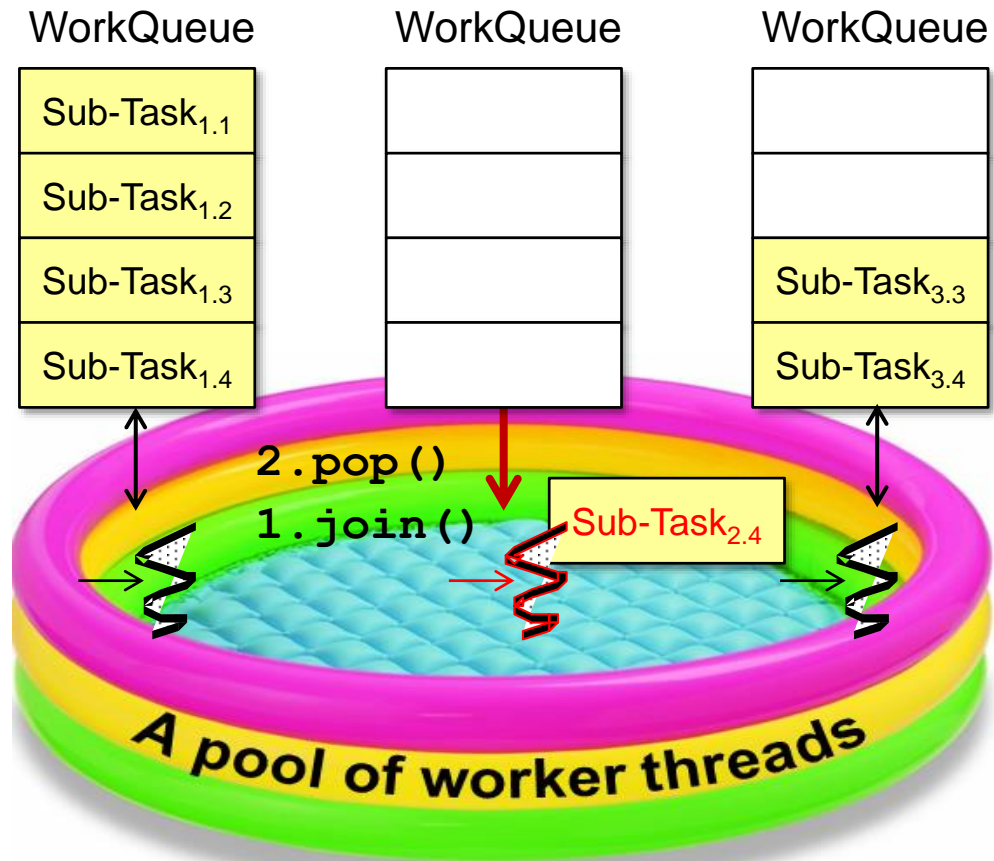
# Java Fork-Join Pool Framework Internals

- If a task run by a worker thread calls `fork()` the new task is pushed on the head of the worker's deque



# Java Fork-Join Pool Framework Internals

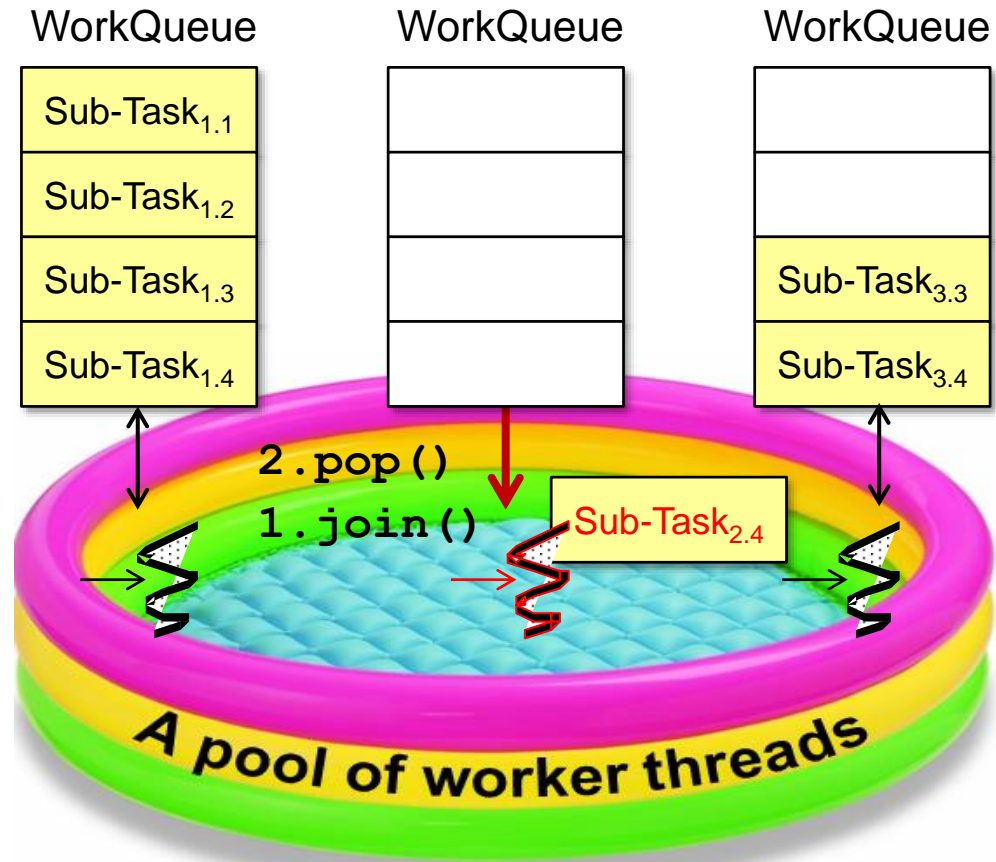
- If a task run by a worker thread calls `fork()` the new task is pushed on the head of the worker's deque
- A worker thread processes its deque in LIFO order



See [en.wikipedia.org/wiki/Stack \(abstract data type\)](https://en.wikipedia.org/wiki/Stack_(abstract_data_type))

# Java Fork-Join Pool Framework Internals

- If a task run by a worker thread calls `fork()` the new task is pushed on the head of the worker's deque
- A worker thread processes its deque in LIFO order, i.e.
  - It pops (sub-)tasks from the head of its deque & runs them to completion

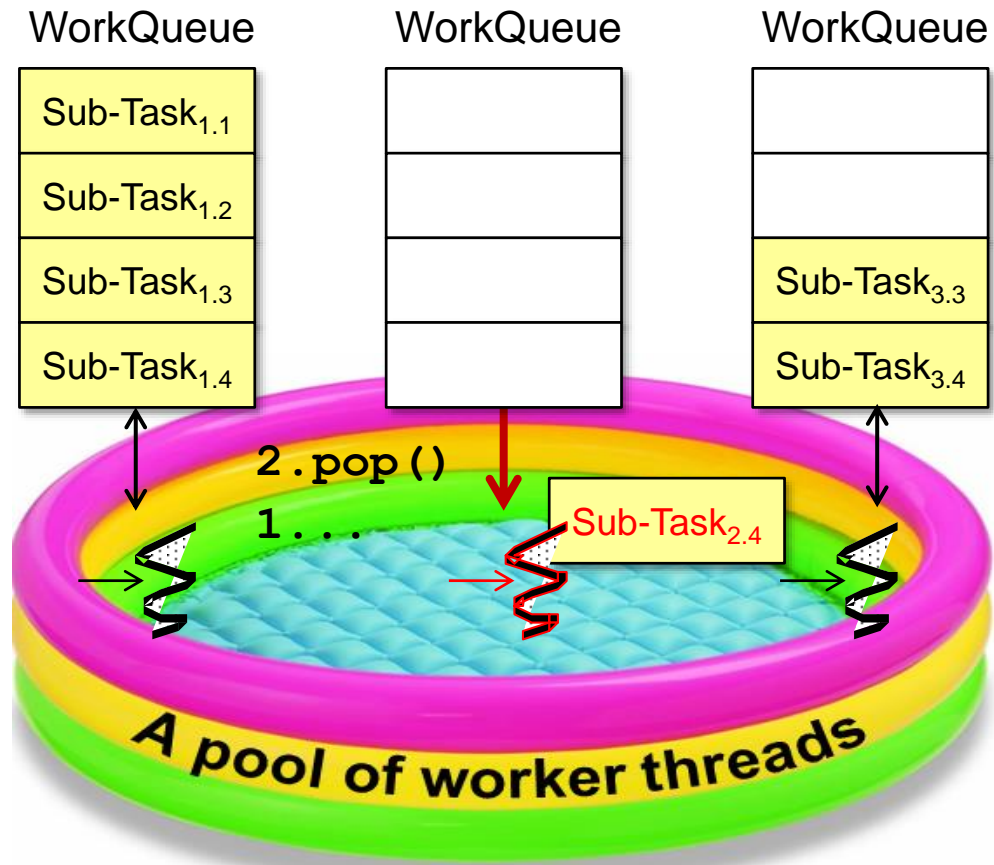


See [en.wikipedia.org/wiki/Run\\_to\\_completion\\_scheduling](https://en.wikipedia.org/wiki/Run_to_completion_scheduling)



# Java Fork-Join Pool Framework Internals

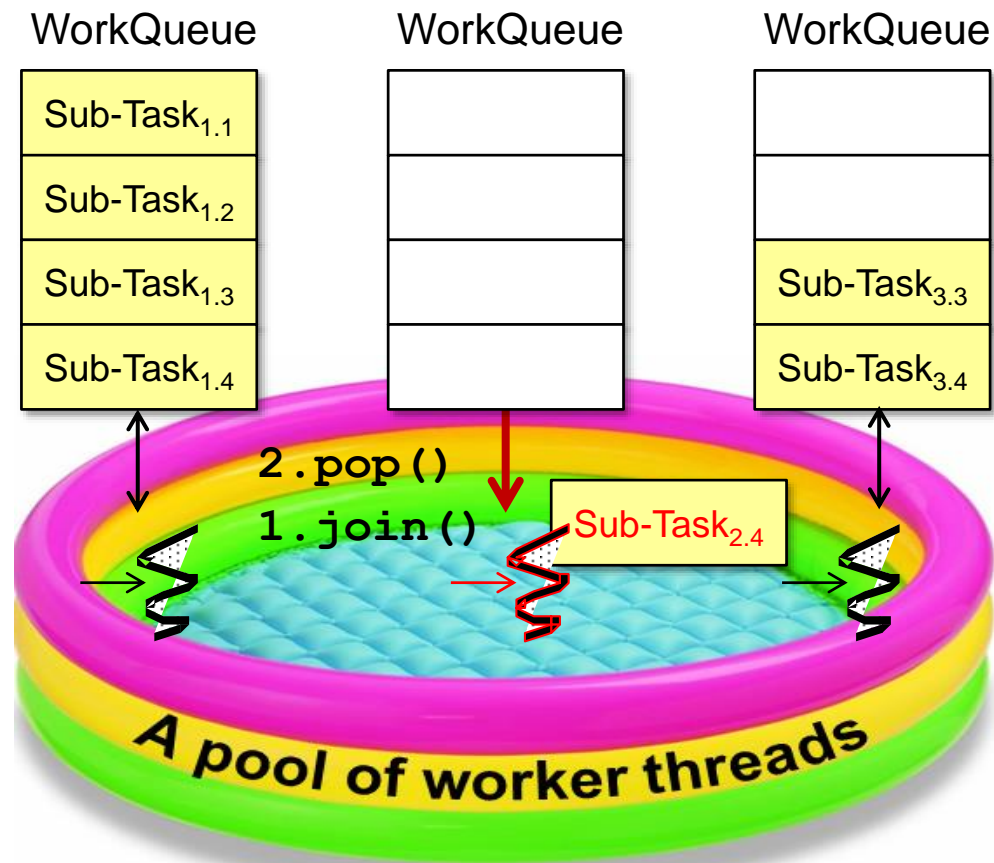
- If a task run by a worker thread calls `fork()` the new task is pushed on the head of the worker's deque
  - A worker thread processes its deque in LIFO order
  - LIFO order improves locality of reference & cache performance



See [en.wikipedia.org/wiki/Locality\\_of\\_reference](https://en.wikipedia.org/wiki/Locality_of_reference)

# Java Fork-Join Pool Framework Internals

- If a task run by a worker thread calls `join()` it pitches in" to pop & execute (sub-)tasks

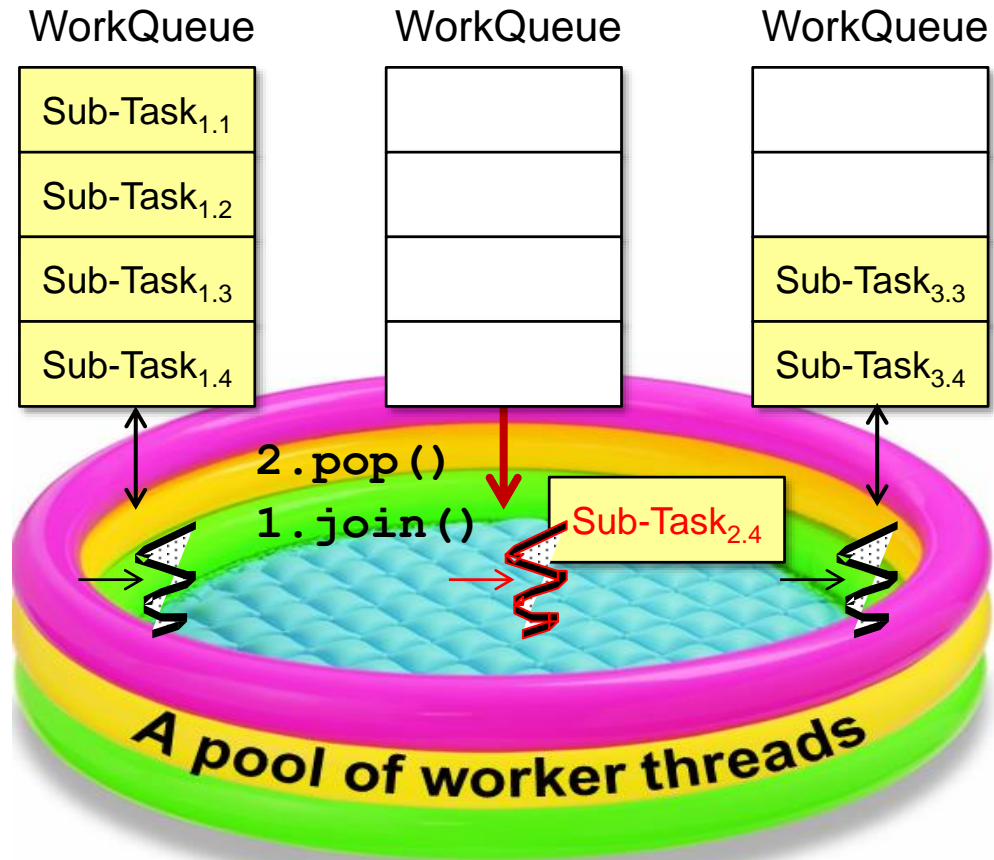


# Java Fork-Join Pool Framework Internals

- If a task run by a worker thread calls `join()` it pitches in" to pop & execute (sub-)tasks




*"Collaborative Jiffy Lube"  
model of processing!*



See [en.wikipedia.org/wiki/Jiffy\\_Lube](https://en.wikipedia.org/wiki/Jiffy_Lube)

# Java Fork-Join Pool Framework Internals


- To maximize core utilization, idle worker threads “steal” work from the tail of busy threads’ dequeues





# Java Fork-Join Pool Framework Internals

- To maximize core utilization, idle worker threads “steal” work from the tail of busy threads’ dequeues





The worker thread deque to steal from is selected randomly to lower contention



# Java Fork-Join Pool Framework Internals


- Tasks are stolen in FIFO order



See [en.wikipedia.org/wiki/FIFO \(computing and electronics\)](https://en.wikipedia.org/wiki/FIFO_(computing_and_electronics))

# Java Fork-Join Pool Framework Internals


- Tasks are stolen in FIFO order, e.g. WorkQueue
- Minimizes contention with thread owning the deque



# Java Fork-Join Pool Framework Internals

- Tasks are stolen in FIFO order, e.g. WorkQueue

- Minimizes contention with thread owning the deque
- An older stolen task may provide a larger unit of work

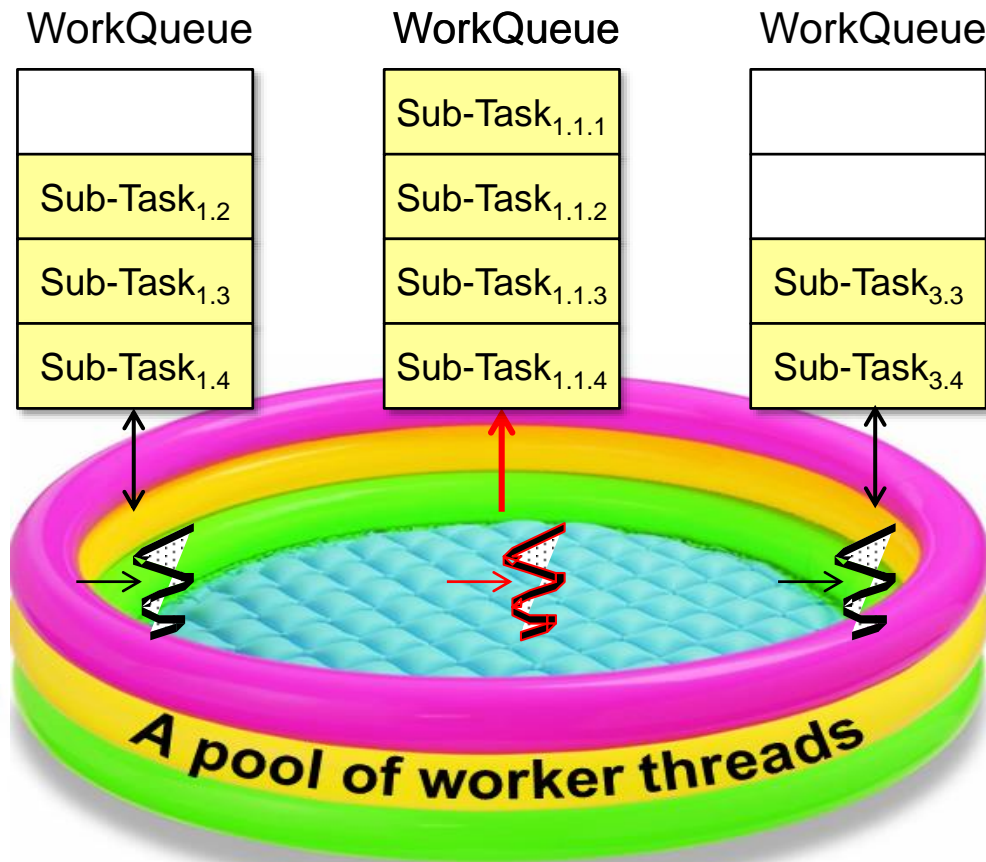




# Java Fork-Join Pool Framework Internals

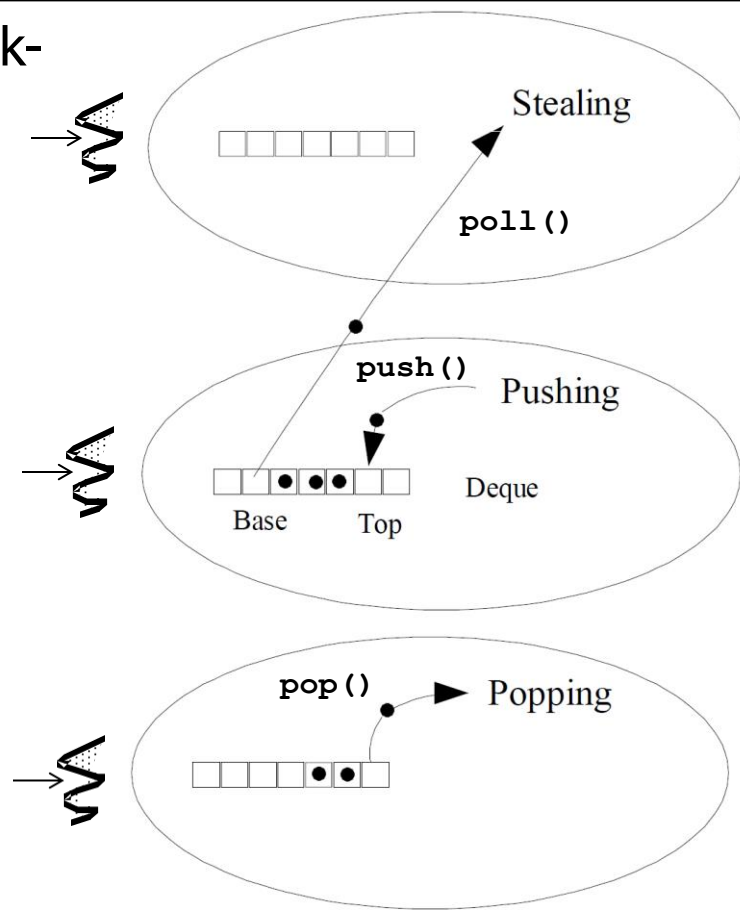
- Tasks are stolen in FIFO order, e.g. WorkQueue

- Minimizes contention with thread owning the deque
- An older stolen task may provide a larger unit of work
  - Enables further recursive decompositions by the stealing thread



# Java Fork-Join Pool Framework Internals

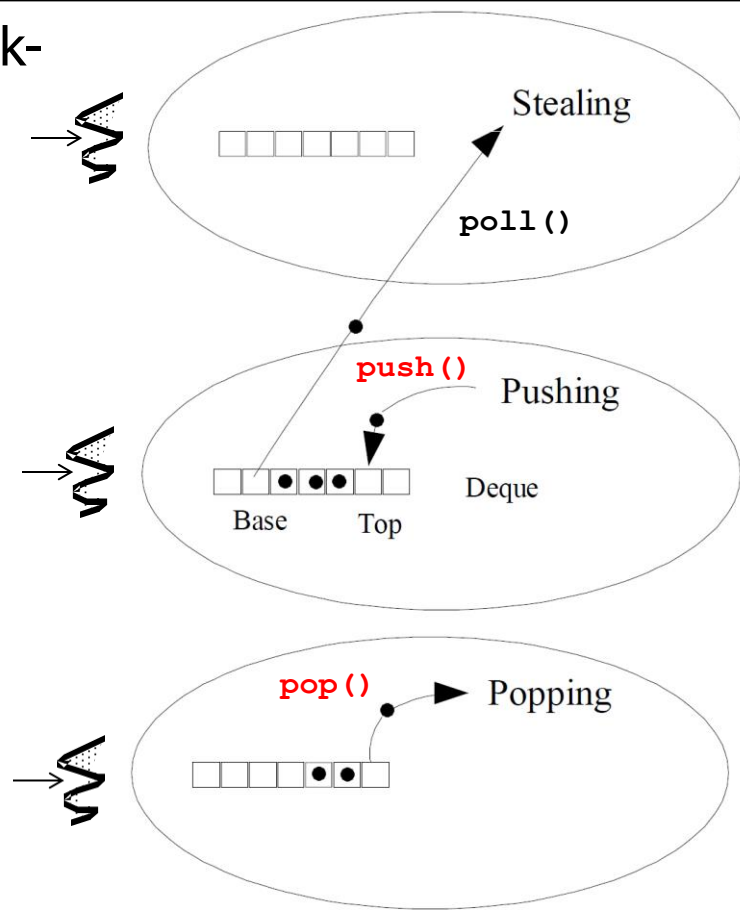
- The WorkQueue deque that implements work-stealing minimizes locking contention





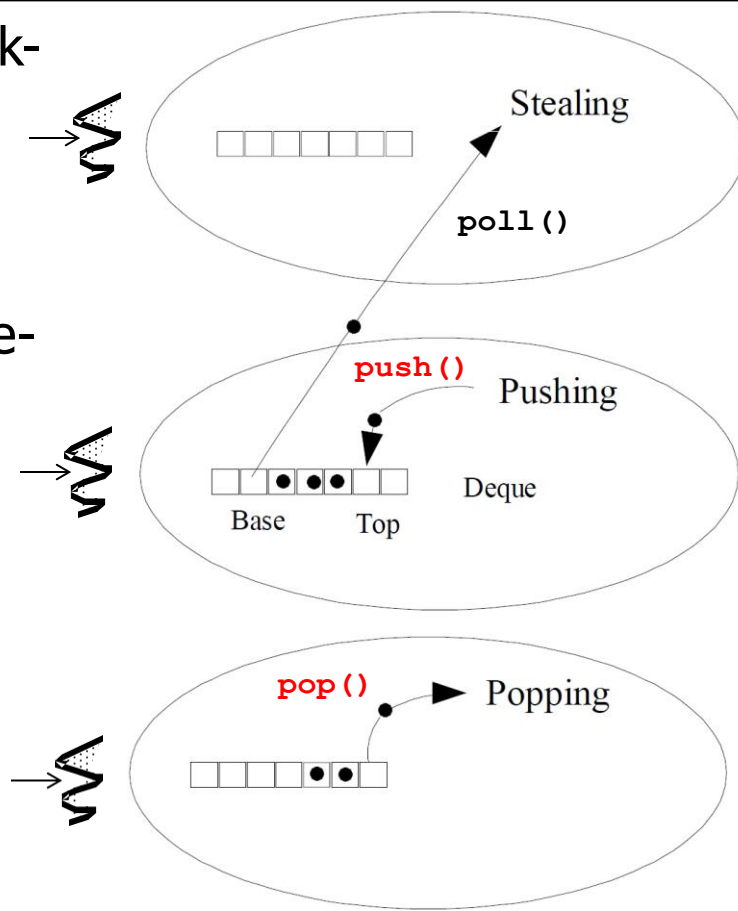
# Java Fork-Join Pool Framework Internals

- The WorkQueue deque that implements work-stealing minimizes locking contention
  - `push()` & `pop()` are only called by the owning worker thread



# Java Fork-Join Pool Framework Internals

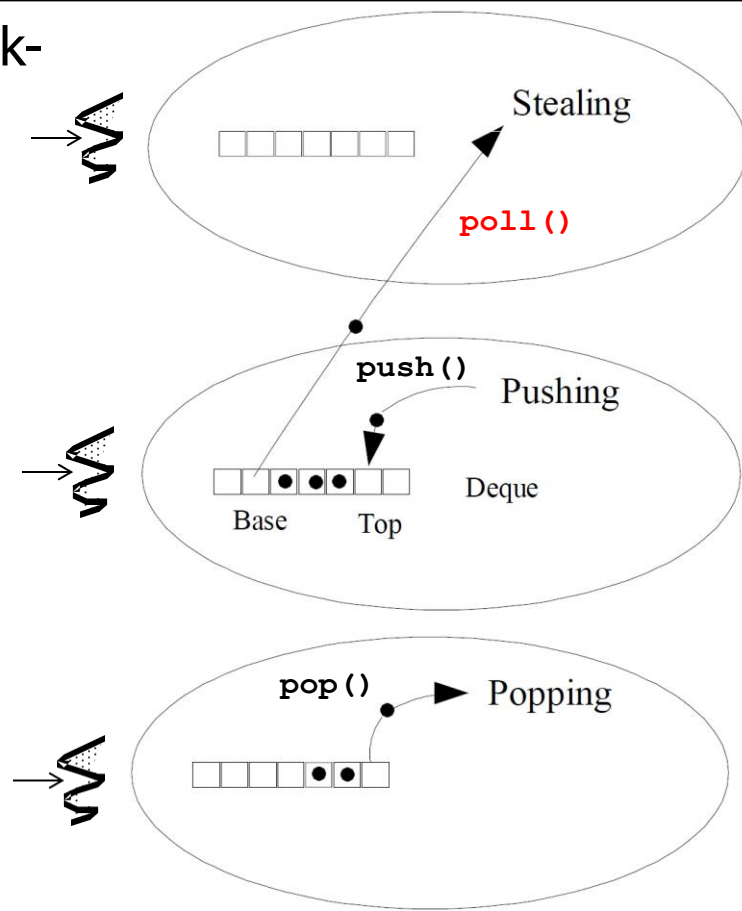
- The WorkQueue deque that implements work-stealing minimizes locking contention
  - `push()` & `pop()` are only called by the owning worker thread
  - These operations use wait-free “compare-and-swap” (CAS) operations



See [en.wikipedia.org/wiki/Compare-and-swap](https://en.wikipedia.org/wiki/Compare-and-swap)

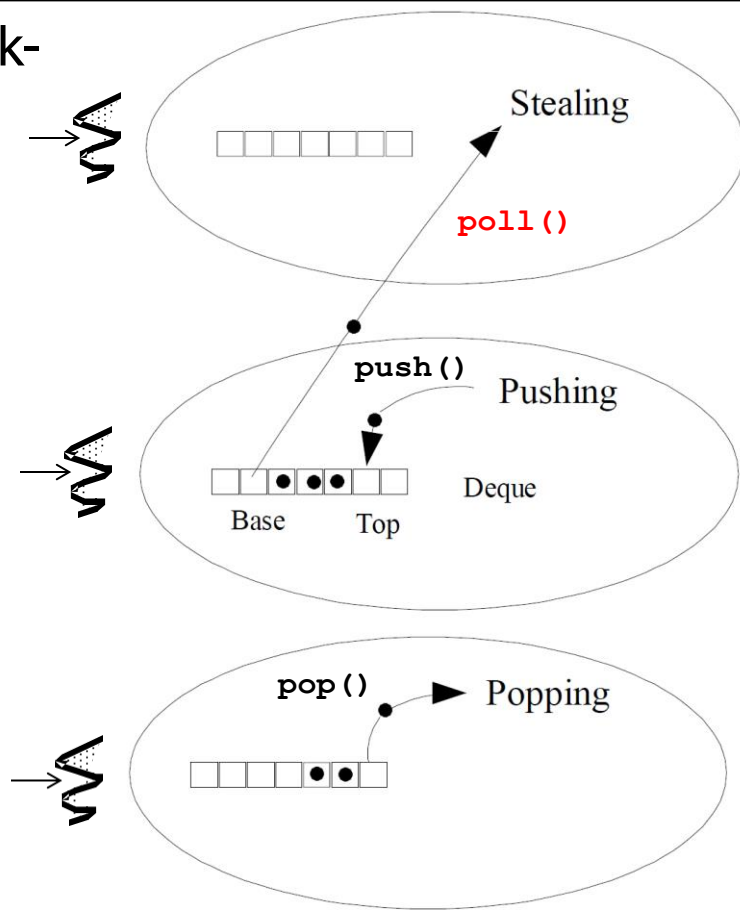
# Java Fork-Join Pool Framework Internals

- The WorkQueue deque that implements work-stealing minimizes locking contention
  - `push()` & `pop()` are only called by the owning worker thread
  - `poll()` may be called from another worker thread to “steal” a (sub-)task



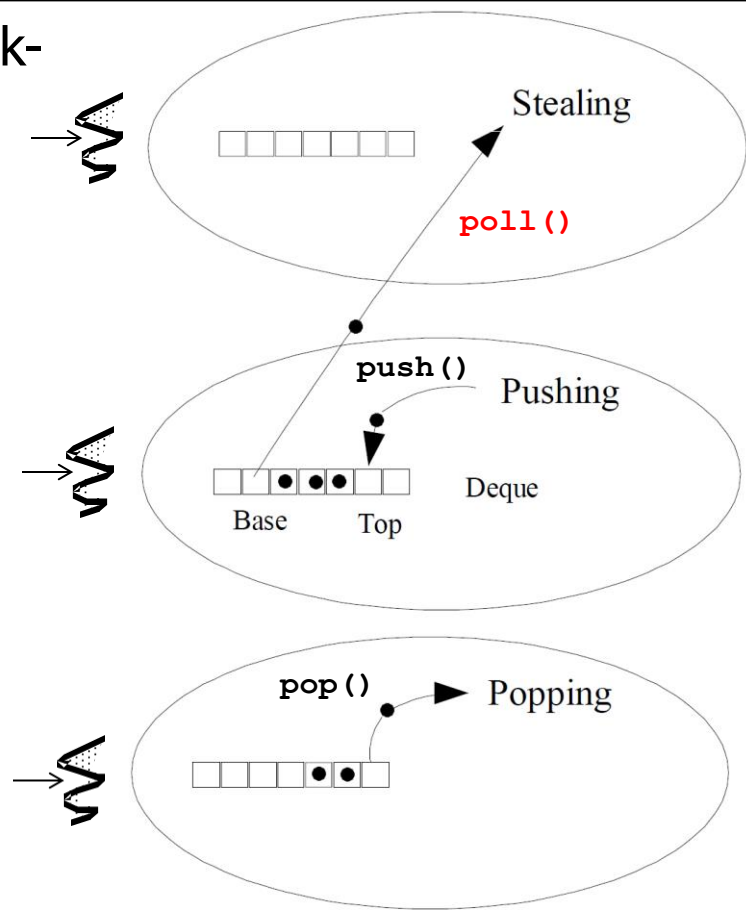
# Java Fork-Join Pool Framework Internals

- The WorkQueue deque that implements work-stealing minimizes locking contention
  - `push()` & `pop()` are only called by the owning worker thread
  - `poll()` may be called from another worker thread to “steal” a (sub-)task
    - May not always be wait-free



# Java Fork-Join Pool Framework Internals

- The WorkQueue deque that implements work-stealing minimizes locking contention
  - `push()` & `pop()` are only called by the owning worker thread
  - `poll()` may be called from another worker thread to “steal” a (sub-)task
  - May not always be wait-free
    - See “Implementation Overview” comments in the `ForkJoinPool` source code for details..



See [java8/util/concurrent/ForkJoinPool.java](https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/ForkJoinPool.java)



---

# End of the Java Fork-Join Pool Framework (Part 3)