

- PG Practise tips:

*Dont do it*

*PG Internal Post*

- check current database's isolation level

```
select current_setting('transaction_isolation');  
-- or  
show transaction_isolation;
```

- Acquire a Snapshot

```
-- get output like: xmin:xmax:list of active transaction ids at  
the time of the snapshot  
-- e.g. 100:104:100,102  
select pg_current_snapshot();
```

- set isolation level within a transaction

```
begin;  
set transaction isolation level repeatable read ;  
show transaction_isolation;  
commit;
```

- How to check pages a table contains

```
-- pg_relpages is in the built-in extension of pgstattuple  
create extension pgstattuple ;  
SELECT pg_relpages('x');
```

- Table Logical Architecture

- pages
- tuples (heap tuple)

- mvcc

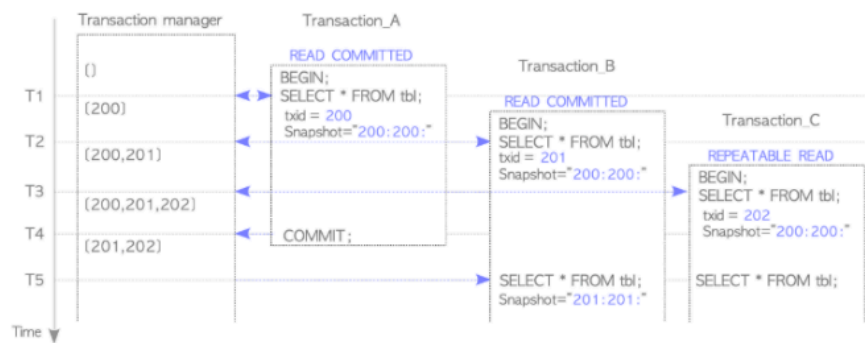
*Due to PostgreSQL's MVCC, when you update a row, PostgreSQL doesn't overwrite the existing tuple.*

*Instead, it creates a new tuple with the updated data*

*so a single logical "row" can have multiple physical "tuples" associated with it, representing its different versions over time*

- SI: snapshot isolation
- Transaction manager: within each transaction, the manager only focus on active transaction
- invisible check: dependent on active transaction and the SI level .
- **almost all tuples within a PostgreSQL table's pages are related to a transaction ID (XID)**

*PostgreSQL's Multi-Version Concurrency Control (MVCC) relies heavily on transaction IDs to manage concurrent data access.*



ree  
COMMITTED,

The transaction manager always holds information about currently running transactions. Suppose that three transactions start one after another, and the isolation level of Transaction\_A and Transaction\_B are READ COMMITTED, and that of Transaction\_C is REPEATABLE READ.

- T1:**  
 Transaction\_A starts and executes the first SELECT command. When executing the first command, Transaction\_A requests the txid and snapshot of this moment.  
 In this scenario, the transaction manager assigns txid 200, and returns the transaction snapshot '200:200:'.
- T2:**  
 Transaction\_B starts and executes the first SELECT command. The transaction manager assigns txid 201, and returns the transaction snapshot '200:200:' because Transaction\_A (txid 200) is in progress. Thus, Transaction\_A cannot be seen from Transaction\_B.
- T3:**  
 Transaction\_C starts and executes the first SELECT command. The transaction manager assigns txid 202, and returns the transaction snapshot '200:200:'. thus, Transaction\_A and Transaction\_B cannot be seen from Transaction\_C.
- T4:**  
 Transaction\_A has been committed. The transaction manager removes the information about this transaction.
- T5:**  
 Transaction\_B and Transaction\_C execute their respective SELECT commands.  
 Transaction\_B requires a transaction snapshot because it is in the READ COMMITTED level.  
 In this scenario, Transaction\_B obtains a new snapshot '201:201:' because Transaction\_A (txid 200) is committed. Thus, Transaction\_A is no longer invisible from Transaction\_B.  
 In contrast, Transaction\_C does not require a transaction snapshot because it is in the REPEATABLE READ level and uses the obtained snapshot, i.e. '200:200:'. Thus, Transaction\_A is still invisible from Transaction\_C.